

# RLM Algorithm: Implementation vs. Paper Comparison

Comparison of our rlm-loop implementation against the original MIT CSAIL paper (arXiv:2512.24601).

---

## Our Algorithm (as implemented)

User Query + Context (file/dir/string)

1. Materialize context → plain Python str
2. Optionally build "context sample" (--no-context-sample disables)
3. Create depth-tracked llm\_query(snippet, task) closure (max\_depth=3)
4. Initialize REPLEnv with namespace: CONTEXT, FILES, llm\_query, FINAL, SHOW\_VARS, pre-imported modules
  
5. Send to LLM: system prompt (full/compact) + user query [+ context sample]

Iteration Loop (max\_iterations, timeout, token budget)

6. Check timeout / token budget guards
7. LLM responds with markdown
8. Extract ```python blocks (regex)
9. If no code blocks:
  - Check for "FINAL" in text → done
  - Else prompt LLM to write code
10. Execute each code block sequentially  
in persistent REPL namespace
11. If FINAL() called → done
12. Feed captured print() output back  
as user message: "Output:\n..."
13. Continue loop

Return RLMResult(answer, stats, history)

Key files: rlm/rlm.py (orchestrator), rlm/repl.py (REPL environment),  
rlm/prompts.py (system prompts), rlm/backends.py (LLM backends)

---

## Paper's Algorithm

User Query + Context

1. Store context as CONTEXT variable in Python REPL
2. Provide `llm_query(snippet, task)` for recursive sub-LM calls
3. Root LM receives ONLY the query (context is not inline)
  
4. Send to LLM: system prompt + query (no context sample)

Code-Observe-Reason Loop

5. LLM generates Python code
6. Execute in REPL sandbox
7. Observe output (truncated)
8. LLM decides next action:
  - More code (inspect/search/chunk)
  - `llm_query()` for semantic analysis
  - `FINAL(answer)` when confident
9. Continue loop

Return final answer

---

## Side-by-Side Comparison

Aspect	Paper	Our Implementation	Delta
<b>Initial context exposure</b>	Root LM sees only the query; CONTEXT is opaque until code runs	Context sample included by default; <code>--no-context-sample</code> disables it	<b>Aligned</b> — can now match paper's opaque-context design via flag
<b>llm_query signature</b>	<code>llm_query(snippet, task)</code> — two args	<code>llm_query(snippet, task) — two args</code>	<b>Match</b>

Aspect	Paper	Our Implementation	Delta
<b>Recursive depth</b>	Depth-1 only (root → sub-LM, no deeper)	Configurable <code>max_depth=3</code> (supports multi-level recursion)	We're ahead of the paper here
<b>Sub-LM model</b>	Smaller/cheaper model for sub-calls (e.g., GPT-5-mini when root is GPT-5)	<code>--recursive-model</code> param exists but defaults to same model as <code>--model</code>	Paper's design is more cost-efficient
<b>System prompt for sub-calls</b>	Not detailed	No system prompt — bare user message only	Consistent but sub-calls might benefit from minimal guidance
<b>Termination</b>	<code>FINAL()</code> + max iterations + timeouts	<code>FINAL()</code> + max iterations + <code>--timeout</code> + <code>--max-token-budget</code>	<b>Match</b>
<b>Default model</b>	Not applicable (paper uses specific models per experiment)	<code>--model</code> is required (no defaults)	<b>Aligned</b> — user must choose explicitly
<b>Code extraction</b>	Not specified in detail	Regex: <code>```python\n(.*)?```</code> — only explicitly tagged Python blocks	Reasonable; ignores non-Python fenced blocks
<b>Sandbox isolation</b>	Rootless container (security delegated to runtime)	Safe builtins whitelist (no <code>__import__</code> , <code>open</code> , <code>exec</code> ) + container delegation	We add in-process restrictions on top of container isolation
<b>Output truncation</b>	Mentioned (first ~10K chars)	10,000 char cap per execution	<b>Match</b>

Aspect	Paper	Our Implementation	Delta
<b>Multi-file support</b>	Not discussed	FILES dict + CompositeContext for directories	Extension beyond the paper
<b>Context loading</b>	Context as string	Memory-mapped files via LazyContext, CompositeContext for multi-file	Engineering refinement beyond the paper
<b>Async execution</b>	Identified as future work	Stubs exist (acompletion) but delegate to sync	Neither implements true async
<b>Native post-training</b>	RLM-Qwen3-8B (28.3% improvement over base)	Prompt-only, no fine-tuning	Paper has a training component we don't
<b>Cost tracking</b>	Discussed (comparable median, high variance tail)	RLMStats tracks tokens, iterations, recursive calls; --max-token-budget enforces limits	<b>Aligned</b> — cost enforcement now available
<b>Verification strategy</b>	Some models perform redundant verification sub-calls	Not prompted for; up to the LLM	Paper observes this as emergent behavior

## Remaining Divergences

### 1. Context Sample in Initial Prompt (our addition, now optional)

**Our approach** (`_build_context_sample()`): By default, sends evenly-spaced 500-char excerpts from the document in the first message. Can be disabled with `--no-context-sample` to match the paper's opaque-context design.

**Paper's approach:** Root LM only knows CONTEXT exists. It must write code to inspect it.

**Trade-off:** The default gives the LLM a head start (knows format, size, structure immediately) but adds tokens to the initial prompt. Use `--no-context-sample` for the paper's strict “context as environment” paradigm.

## 2. Multi-level Recursion (our extension)

**Paper:** Depth-1 only (root → sub-LM, no deeper).

**Ours:** Configurable `max_depth=3` (supports multi-level recursion).

This is an intentional extension beyond the paper’s design.

---

## Emergent Strategy Comparison

Both the paper and our prompts describe the same four-phase strategy:

Phase	Paper	Our Prompts
<b>Inspect</b>	Examine context structure, size, format	“Study the document sample... understand format”
<b>Search</b>	Regex/keyword filtering to narrow focus	“Start BROAD: <code>re.findall(pattern, CONTEXT)</code> with <code>re.IGNORECASE</code> ”
<b>Chunk</b>	Uniform or semantic partitioning	“Extract sections: <code>CONTEXT[start:end]</code> (aim for 1000-5000 chars)”
<b>Synthesize</b>	Aggregate sub-LM results + <code>FINAL()</code>	“Combine and call <code>FINAL(answer)</code> ”

Our prompts are more prescriptive (explicit chunk size guidance, specific function examples). The paper observes these strategies as emergent behavior that models discover with minimal prompting.

---

## Future Refinement Candidates

Actionable items identified by this comparison:

- **Default to a cheaper model for recursive sub-calls** — leverage `--recursive-model` param with a smaller default
- **Evaluate prompt prescriptiveness** — test whether our detailed prompts help or constrain the LLM compared to the paper’s minimal approach