

# RLM Algorithm: Implementation vs. Paper Comparison

Comparison of our `rlm-loop` implementation against the original MIT CSAIL paper (arXiv:2512.24601).

---

## Our Algorithm (as implemented)

User Query + Context (file/dir/string)

1. Materialize context → plain Python str
2. Optionally build "context sample" (`--no-context-sample` disables)
3. Create depth-tracked `llm_query(snippet, task)` closure (`max_depth=3`)
4. Initialize REPLEnv with namespace: CONTEXT, FILES, `llm_query`, FINAL, SHOW\_VARS, pre-imported modules
  
5. Send to LLM: system prompt (full/compact) + user query [+ context sample]

RLM Loop (max\_iterations, timeout, token budget)

6. Check timeout / token budget guards
7. LLM responds with markdown
8. Extract ```python blocks (regex)
9. If no code blocks:
  - Check for "FINAL" in text → done
  - Else prompt LLM to write code
10. Execute each code block sequentially  
in persistent REPL namespace
11. If `FINAL()` called → done
12. Feed captured `print()` output back  
as user message: "Output:\n..."
13. Continue loop

Return `RLMResult(answer, stats, history)`

Key files: `rlm/rlm.py` (orchestrator), `rlm/repl.py` (REPL environment),  
`rlm/prompts.py` (system prompts), `rlm/backends.py` (LLM backends)

---

## Paper's Algorithm

User Query + Context

1. Store context as CONTEXT variable in Python REPL
2. Provide `llm_query(snippet, task)` for sub-RLM calls
3. Root LM receives ONLY the query (context is not inline)
  
4. Send to LLM: system prompt + query (no context sample)

Code-Observe-Reason Loop

5. LLM generates Python code
6. Execute in REPL sandbox
7. Observe output (truncated)
8. LLM decides next action:
  - More code (`inspect/search/chunk`)
  - `llm_query()` for semantic analysis
  - `FINAL(answer)` when confident
9. Continue loop

Return final answer

---

## Side-by-Side Comparison

Aspect	Paper	Our Implementation	Delta
<b>Initial context exposure</b>	Root LM sees only the query; CONTEXT is opaque until code runs	Context sample included by default; <code>--no-context-sample</code> disables it	<b>Aligned</b> — can now match paper's opaque-context design via flag
<b>llm_query signature</b>	<code>llm_query(snippet, task)</code> — two args	<code>llm_query(snippet, task) — two args</code>	<b>Match</b>

Aspect	Paper	Our Implementation	Delta
<b>Recursive depth</b>	Depth-1 only (root → sub-RLM, no deeper)	Configurable <code>max_depth=3</code> (supports multi-level recursion)	We're ahead of the paper here
<b>Sub-RLM model</b>	Smaller/cheaper <code>sub-rlm-model</code> param exists model for sub-calls (e.g., GPT-5-mini when root is GPT-5)	but defaults to same model as <code>--model</code>	Paper's design is more cost-efficient
<b>System prompt for sub-calls</b>	Not detailed	Minimal task-focused system prompt (text analysis role)	Aligned — sub-calls now receive focused guidance
<b>Termination</b>	<code>FINAL()</code> + max iterations + timeouts	<code>FINAL()</code> + max iterations + <code>--timeout</code> + <code>--max-token-budget</code>	<b>Match</b>
<b>Default model</b>	Not applicable (paper uses specific models per experiment)	<code>--model</code> is required (no defaults)	<b>Aligned</b> — user must choose explicitly
<b>Code extraction</b>	Not specified in detail	Regex: <code>```python\n(.*?)```</code> — only explicitly tagged Python blocks	Reasonable; ignores non-Python fenced blocks
<b>Sandbox isolation</b>	Rootless container (security delegated to runtime)	Safe builtins whitelist (no <code>__import__</code> , <code>open</code> , <code>exec</code> ) + container delegation	We add in-process restrictions on top of container isolation
<b>Output truncation</b>	Mentioned (first ~10K chars)	10,000 char cap per execution	Match

Aspect	Paper	Our Implementation	Delta
<b>Multi-file support</b>	Not discussed	FILES dict + CompositeContext for directories	Extension beyond the paper
<b>Context loading</b>	Context as string	Memory-mapped files via LazyContext, CompositeContext for multi-file	Engineering refinement beyond the paper
<b>Async execution</b>	Identified as future work	Stubs exist (acompletion) but delegate to sync	Neither implements true async
<b>Native post-training</b>	RLM-Qwen3-8B (28.3% improvement over base)	Prompt-only, no fine-tuning	Paper has a training component we don't
<b>Cost tracking</b>	Discussed (comparable median, high variance tail)	RLMStats tracks tokens, iterations, sub-RLM calls; --max-token-budget enforces limits	<b>Aligned</b> — cost enforcement now available
<b>Verification strategy</b>	Some models perform redundant verification sub-calls	--verify flag enables explicit verification sub-call on the final answer	Paper observes this as emergent behavior; we formalize it as opt-in

## LLM Roles

The RLM pattern uses two distinct LLM tiers with different responsibilities:

### Role Comparison

Aspect	Root LM	Sub-RLM
Depth	0	1+ (up to <code>max_depth</code> )

System prompt Full strategy prompt (~120 lines) or compact (~25 lines)

Aspect	Root LM	Sub-RLM
Conversation	Multi-turn (code → output → code → ... → FINAL)	Single-turn (one prompt, one response)
Produces	Python code in ```python blocks	Plain text (summary, extraction, analysis)
Has access to	CONTEXT, FILES, llm_query(), FINAL(), SHOW_VARS(), pre-imported modules	Only the snippet passed by root
Model	--model	--sub-rlm-model (defaults to --model)

### Root LM Responsibilities

The root LM writes Python code to explore the document programmatically. It never sees the full document inline (unless `--no-context-sample` is disabled). Instead, it:

1. Receives the user query (and optionally a document sample)
2. Writes exploration code to inspect structure, search for patterns, extract chunks
3. Delegates **semantic** analysis to the sub-RLM via `llm_query(snippet, task)`
4. Aggregates sub-RLM results and calls `FINAL(answer)` when confident

### Sub-RLM Responsibilities

The sub-RLM processes individual text chunks. It:

1. Receives a snippet of text and a task description
2. Returns plain-text results (summaries, extractions, classifications)
3. Has no access to the REPL, CONTEXT, or any tools — only the snippet it was given
4. Does not generate code

### Emergent Strategies from the Paper

The paper (arXiv:2512.24601) and the accompanying blog post identify five emergent strategies that models discover when given the RLM environment:

Strategy	Description	Example Code
<b>Peeking</b>	Inspect document structure and size	<code>print(CONTEXT[:1000]), print(len(CONTEXT))</code>
<b>Grepping</b>	Keyword/regex search to locate relevant regions	<code>re.findall(r'pattern', CONTEXT, re.IGNORECASE)</code>

Strategy	Description	Example Code
<b>Partition+Map</b>	Chunk the document and fan out <code>llm_query()</code> calls	Split <code>CONTEXT</code> into $N$ chunks, call <code>llm_query()</code> on each
<b>Summarization</b>	Hierarchical summarization via sub-RLM calls	Summarize chunks, then summarize the summaries
<b>Long-input/long-output</b>	Direct extraction for structured outputs	Extract lists, tables, or formatted data from large regions

Our prompts map these emergent strategies to a prescriptive four-phase model:

Our Phase	Paper Strategies Covered
<b>Inspect</b>	Peeking
<b>Search</b>	Grepping
<b>Chunk + Analyze</b>	Partition+Map, Summarization
<b>Synthesize</b>	Long-input/long-output, final aggregation

The paper observes that models discover these strategies naturally with minimal prompting. Our prompts are more prescriptive (explicit chunk size guidance, specific function examples), which improves consistency but may constrain creative exploration.

### Verification (Emergent Behavior)

The paper observes that some models perform redundant “verification” sub-calls — re-checking their own intermediate results before calling `FINAL()`. This behavior is emergent and not prompted for. Our implementation offers an explicit `--verify` flag that formalizes this pattern as an opt-in post-processing step.

---

## Remaining Divergences

### 1. Context Sample in Initial Prompt (our addition, now optional)

**Our approach (`_build_context_sample()`):** By default, sends evenly-spaced 500-char excerpts from the document in the first message. Can be disabled with `--no-context-sample` to match the paper’s opaque-context design.

**Paper’s approach:** Root LM only knows `CONTEXT` exists. It must write code to inspect it.

**Trade-off:** The default gives the LLM a head start (knows format, size, structure immediately) but adds tokens to the initial prompt. Use `--no-context-sample` for the paper’s strict “context as environment” paradigm.

## 2. Multi-level Recursion (our extension)

**Paper:** Depth-1 only (root → sub-RLM, no deeper).

**Ours:** Configurable `max_depth=3` (supports multi-level recursion).

This is an intentional extension beyond the paper’s design.

---

## Emergent Strategy Comparison

Both the paper and our prompts describe the same four-phase strategy:

Phase	Paper	Our Prompts
<b>Inspect</b>	Examine context structure, size, format	If sample provided, study it; otherwise <code>print(CONTEXT[:1000])</code>
<b>Search</b>	Regex/keyword filtering to narrow focus	“Start BROAD: <code>re.findall(pattern, CONTEXT)</code> with <code>re.IGNORECASE</code> ”
<b>Chunk</b>	Uniform or semantic partitioning	“Extract sections: <code>CONTEXT[start:end]</code> (aim for 1000-5000 chars)”
<b>Synthesize</b>	Aggregate sub-RLM results + <code>FINAL()</code>	“Combine and call <code>FINAL(answer)</code> ”

Our prompts also include guidance on when to use `llm_query()` (semantic tasks) vs Python directly (structural tasks). The paper observes these strategies as emergent behavior that models discover with minimal prompting; our prompts are more prescriptive to improve consistency.

---

## Future Refinement Candidates

Actionable items identified by this comparison:

- **Default to a cheaper model for sub-RLM calls** — leverage `--sub-rlm-model` param with a smaller default

- **Evaluate prompt prescriptiveness** — test whether our detailed prompts help or constrain the LLM compared to the paper’s minimal approach
- **Native post-training** — the paper fine-tuned RLM-Qwen3-8B (28.3% improvement over base); we rely on prompt-only guidance
- **True async execution** — replace sync-delegating acompletion stubs with genuine async REPL execution